



Data Cleaning with Python

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
import warnings
warnings.filterwarnings('ignore')
import re
```

Data Collection and Inspection

```
In [2]: # file reading
df = pd.read_csv('messy_customer_sales_data.csv')
df
```

	Customer_ID	Name	Gender	Age	City	Signup_Date	Last_Pur
0	CUST4371	Paul Wilson	m	52.0	KOLKATA	2025-06-26	
1	CUST5957	Jason Thomas	M	51.0 years	NaN	2021-02-17	
2	CUST3754	Brittney Martinez	F	62.0	hyderabad	2023-11-05	
3	CUST2934	Brenda Pierce	FEMALE	40.0	hyderabad	2022-03-13	
4	CUST5683	Matthew Carroll	f	41.0	CHENNAI	2024-04-05	
...
10195	CUST10767	Robert Lewis	female	35.0 years	delhi	2020-12-08	
10196	NaN	Diane Evans	M	53.0	bangalore	2023-12-31	
10197	CUST6315	Joshua Martinez	m	25.0	hyderabad	2022-02-15	
10198	CUST4812	Sarah Miller	FEMALE	55.0	NaN	2021-03-16	
10199	CUST6588	David Potter	female	34.0	HYDERABAD	2020-10-12	

10200 rows × 12 columns

```
In [3]: # data types  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10200 entries, 0 to 10199  
Data columns (total 12 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   Customer_ID     9177 non-null    object    
 1   Name             10200 non-null   object    
 2   Gender           9174 non-null    object    
 3   Age              9249 non-null    object    
 4   City             9184 non-null    object    
 5   Signup_Date      10200 non-null   object    
 6   Last_Purchase_Date 9188 non-null   object    
 7   Purchase_Amount  9179 non-null    float64   
 8   Feedback_Score   9177 non-null    float64   
 9   Email            10200 non-null   object    
 10  Phone_Number     10200 non-null   int64     
 11  Country          9468 non-null    object    
dtypes: float64(2), int64(1), object(9)  
memory usage: 956.4+ KB
```

```
In [4]: # summary statistics  
df.describe()
```

	Purchase_Amount	Feedback_Score	Phone_Number
count	9.179000e+03	9177.000000	1.020000e+04
mean	2.909013e+04	5.479351	4.979974e+09
std	2.086971e+05	2.867123	2.902593e+09
min	-5.000000e+02	1.000000	9.208990e+05
25%	1.229500e+04	3.000000	2.449157e+09
50%	2.433000e+04	5.000000	4.988639e+09
75%	3.713000e+04	8.000000	7.510448e+09
max	9.999999e+06	10.000000	9.994402e+09

```
In [5]: # missing value summary  
df.isnull().sum()
```

```
Out[5]: Customer_ID      1023
         Name            0
         Gender        1026
         Age           951
         City          1016
         Signup_Date     0
         Last_Purchase_Date 1012
         Purchase_Amount 1021
         Feedback_Score 1023
         Email          0
         Phone_Number    0
         Country        732
         dtype: int64
```

```
In [6]: # total number of row & columns
df.shape
```

```
Out[6]: (10200, 12)
```

```
In [7]: df.dropna().shape
```

```
Out[7]: (4528, 12)
```

```
In [8]: # duplicate rows summary
df[df.duplicated()].shape
```

```
Out[8]: (15, 12)
```

```
In [9]: df[df.duplicated()]
```

Out[9] :

	Customer_ID	Name	Gender	Age	City	Signup_Date	Last_Purchase
2032	CUST2755	Casey Campbell	male	NaN	Hyderabad	2023-12-01	
3191	NaN	Travis Schneider	f	51.0	NaN	2025-04-20	
4312	CUST8305	Stanley Cain	FEMALE	51.0	Bangalore	2023-08-14	
5611	CUST6288	Jonathon Kim	male	46.0	Chennai	2023-11-22	
5680	CUST2695	Lisa Durham	NaN	46.0	Chennai	2021-04-18	
5899	CUST8841	Amanda Hill	NaN	34.0	chennai	2022-09-11	
6070	CUST10780	Emily Smith	NaN	23.0	Kolkata	2023-12-09	
6240	CUST9745	David Morales	NaN	56.0	NaN	2023-02-18	
6425	CUST4631	George Villa	m	26.0	Mumbai	2022-05-21	
6498	CUST1711	Stephanie Elliott	NaN	44.0	MUMBAI	2023-05-17	
7118	CUST9380	Paul Wilson	MALE	42.0	Bangalore	2025-07-31	
8392	CUST1436	Ralph Anderson	female	33.0	Kolkata	2022-10-21	
9155	CUST10040	Mark Taylor	NaN	43.0	Delhi	2021-04-16	
9960	CUST6220	John Rodriguez	male	51.0	Hyderabad	2021-03-28	
10140	CUST4367	Amber Kennedy	F	61.0	Bangalore	2023-11-05	

In [10] :

`df['Customer_ID'].value_counts()`

```
Out[10]: Customer_ID
CUST3344      2
CUST4893      2
CUST7000      2
CUST10824     2
CUST2695      2
..
CUST7857      1
CUST3881      1
CUST1565      1
CUST9038      1
CUST5957      1
Name: count, Length: 9000, dtype: int64
```

```
In [11]: df[df['Customer_ID'] == 'CUST7000']
```

	Customer_ID	Name	Gender	Age	City	Signup_Date	Last_Purchase_Da
1508	CUST7000	Brittany Ortiz	FEMALE	NaN	Delhi	2022-03-27	2025-02-
3194	CUST7000	Brittany Ortiz	f	NaN	Delhi	2022-03-27	2025-02-

```
In [12]: # category count
for col in df.columns:
    if df[col].nunique() < 20:
        print(df[col].value_counts())
    print('-'*50)
```

```
Gender
f          1184
M          1171
m          1163
F          1157
MALE       1131
female     1128
male       1121
FEMALE     1119
Name: count, dtype: int64
```

```
-----
City
Kolkata      820
Mumbai        812
Chennai       784
Bangalore     773
Hyderabad    770
Delhi         763
CHENNAI       404
KOLKATA       395
MUMBAI        393
hyderabad    384
bangalore     383
DELHI         378
delhi         369
BANGALORE     363
HYDERABAD     360
mumbai        352
chennai        343
kolkata       338
Name: count, dtype: int64
```

```
-----
Feedback_Score
2.0          952
4.0          947
7.0          938
6.0          927
3.0          913
8.0          912
1.0          907
9.0          903
10.0         901
5.0          877
Name: count, dtype: int64
```

```
-----
Country
India        7132
IND          793
india        772
InDia        771
Name: count, dtype: int64
```

Handling missing data

```
In [13]: # Drop rows with missing 'Customer_ID' (unique identifier)
df.dropna(subset = ['Customer_ID'], inplace = True)
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: Customer_ID      0
Name          0
Gender       934
Age         859
City        918
Signup_Date     0
Last_Purchase_Date 914
Purchase_Amount   927
Feedback_Score    905
Email          0
Phone_Number     0
Country       664
dtype: int64
```

```
In [15]: df['Age'].unique()
```

```
Out[15]: array(['52.0', '51.0 years', '62.0', '40.0', '41.0', nan, '18.0',
       '43.0 years', '40.0 years', '26.0', '32.0', '22.0', '59.0', '65.0',
       '61.0', '31.0', '54.0 years', '55.0', '69.0', '61.0 years', '24.0',
       '63.0', '19.0', '50.0', '56.0', '36.0', '68.0', '43.0', '38.0',
       '27.0', '57.0 years', '23.0', '25.0', '66.0', '28.0', '30.0',
       '46.0', '48.0', '20.0', '37.0', '67.0', '51.0', '35.0', '58.0',
       '29.0', 'nan years', '39.0', '49.0', '47.0', '42.0', '44.0',
       '64.0', '53.0', '60.0', '59.0 years', '45.0', '21.0', '34.0',
       '54.0', '48.0 years', '46.0 years', '33.0', '57.0', '30.0 years',
       '58.0 years', '35.0 years', '34.0 years', '69.0 years', '250',
       '19.0 years', '27.0 years', '53.0 years', '65.0 years',
       '66.0 years', '44.0 years', '49.0 years', '25.0 years',
       '23.0 years', '62.0 years', '41.0 years', '33.0 years',
       '28.0 years', '22.0 years', '20.0 years', '42.0 years',
       '45.0 years', '3', '63.0 years', '37.0 years', '38.0 years',
       '55.0 years', '18.0 years', '36.0 years', '67.0 years',
       '29.0 years', '39.0 years', '31.0 years', '47.0 years',
       '52.0 years', '-10', '60.0 years', '24.0 years', '26.0 years',
       '21.0 years', '64.0 years', '50.0 years', '68.0 years',
       '32.0 years', '56.0 years'], dtype=object)
```

```
In [16]: def extract_age(age):
    age_num = re.findall('[0-9]+', str(age))
    if len(age_num) > 0:
        return age_num[0]
    else:
        return age
```

```
df['Age'] = df['Age'].apply(lambda x: extract_age(x))

In [17]: df_age = df[df['Age'] != 'nan years']['Age']

In [18]: age_median = int(df_age.dropna().astype('int64').median())

In [19]: age_median

Out[19]: 43

In [20]: # replace nan with median age and extracting numbers
df['Age'].replace('nan years', age_median, inplace = True)

In [21]: df['Age'].fillna(age_median, inplace = True)

In [22]: df['Age'].unique()

Out[22]: array(['52', '51', '62', '40', '41', '43', '18', '43', '26', '32', '22',
       '59', '65', '61', '31', '54', '55', '69', '24', '63', '19', '50',
       '56', '36', '68', '38', '27', '57', '23', '25', '66', '28', '30',
       '46', '48', '20', '37', '67', '35', '58', '29', '39', '49', '47',
       '42', '44', '64', '53', '60', '45', '21', '34', '33', '250', '3',
       '10'], dtype=object)

In [23]: df['Purchase_Amount'].fillna(df['Purchase_Amount'].median(), inplace = True)

In [24]: df['Feedback_Score'].fillna(df['Feedback_Score'].mode()[0], inplace = True)

In [25]: df.isnull().sum()

Out[25]: Customer_ID          0
         Name                0
         Gender             934
         Age                0
         City               918
         Signup_Date        0
         Last_Purchase_Date 914
         Purchase_Amount     0
         Feedback_Score      0
         Email               0
         Phone_Number        0
         Country              664
         dtype: int64

In [26]: for col in ['Gender', 'City', 'Country']:
         df[col].fillna(df[col].mode()[0], inplace = True)

In [27]: df['Last_Purchase_Date'].ffill(inplace = True)
```

Fixing Inconsistent Formatting

```
In [28]: df['Gender'].unique()
```

```
Out[28]: array(['m ', 'M', 'F', 'FEMALE', 'f ', 'male', 'MALE', 'female'],
              dtype=object)
```

```
In [29]: df['Gender'] = df['Gender'].str.strip().str.lower()
```

```
In [30]: df['Gender'].replace({'m':'male','f':'female'}, inplace = True)
```

```
In [31]: df['City'] = df['City'].str.lower().str.strip()
```

```
In [32]: df['Country'] = df['Country'].str.lower()
```

```
In [33]: df['Country'].replace({'ind':'india'}, inplace = True)
```

```
In [34]: df
```

	Customer_ID	Name	Gender	Age	City	Signup_Date	Last_Purc
0	CUST4371	Paul Wilson	male	52	kolkata	2025-06-26	
1	CUST5957	Jason Thomas	male	51	kolkata	2021-02-17	
2	CUST3754	Brittney Martinez	female	62	hyderabad	2023-11-05	
3	CUST2934	Brenda Pierce	female	40	hyderabad	2022-03-13	
4	CUST5683	Matthew Carroll	female	41	chennai	2024-04-05	
...
10194	CUST6146	Cody Thompson	female	40	kolkata	2024-08-21	
10195	CUST10767	Robert Lewis	female	35	delhi	2020-12-08	
10197	CUST6315	Joshua Martinez	male	25	hyderabad	2022-02-15	
10198	CUST4812	Sarah Miller	female	55	kolkata	2021-03-16	
10199	CUST6588	David Potter	female	34	hyderabad	2020-10-12	

9177 rows × 12 columns

Handling Duplicates

```
In [35]: df[df.duplicated()].shape
```

```
Out[35]: (163, 12)
```

```
In [36]: df.drop_duplicates(inplace = True)
```

```
In [37]: df
```

Out[37]:

	Customer_ID	Name	Gender	Age	City	Signup_Date	Last_Purc
0	CUST4371	Paul Wilson	male	52	kolkata	2025-06-26	
1	CUST5957	Jason Thomas	male	51	kolkata	2021-02-17	
2	CUST3754	Brittney Martinez	female	62	hyderabad	2023-11-05	
3	CUST2934	Brenda Pierce	female	40	hyderabad	2022-03-13	
4	CUST5683	Matthew Carroll	female	41	chennai	2024-04-05	
...
10193	CUST6352	Isaiah Terry	female	26	kolkata	2023-03-04	
10194	CUST6146	Cody Thompson	female	40	kolkata	2024-08-21	
10195	CUST10767	Robert Lewis	female	35	delhi	2020-12-08	
10197	CUST6315	Joshua Martinez	male	25	hyderabad	2022-02-15	
10198	CUST4812	Sarah Miller	female	55	kolkata	2021-03-16	

9014 rows × 12 columns

```
In [38]: df['Customer_ID'].value_counts()
```

```
Out[38]: Customer_ID
CUST3693    2
CUST6833    2
CUST5341    2
CUST1002    2
CUST9950    2
...
CUST7857    1
CUST3881    1
CUST1565    1
CUST9038    1
CUST5957    1
Name: count, Length: 9000, dtype: int64
```

```
In [39]: df[df['Customer_ID'] == 'CUST3693']
```

	Customer_ID	Name	Gender	Age	City	Signup_Date	Last_Purcha
3967	CUST3693	Chad Dominguez	female	3	chennai	2022-01-20	2022-01-20
4953	CUST3693	Chad Dominguez	female	54	chennai	2022-01-20	2022-01-20

```
In [40]: df.drop_duplicates(subset = ['Customer_ID'], keep = 'first')
```

```
Out[40]:
```

	Customer_ID	Name	Gender	Age	City	Signup_Date	Last_Purc
0	CUST4371	Paul Wilson	male	52	kolkata	2025-06-26	
1	CUST5957	Jason Thomas	male	51	kolkata	2021-02-17	
2	CUST3754	Brittney Martinez	female	62	hyderabad	2023-11-05	
3	CUST2934	Brenda Pierce	female	40	hyderabad	2022-03-13	
4	CUST5683	Matthew Carroll	female	41	chennai	2024-04-05	
...
10193	CUST6352	Isaiah Terry	female	26	kolkata	2023-03-04	
10194	CUST6146	Cody Thompson	female	40	kolkata	2024-08-21	
10195	CUST10767	Robert Lewis	female	35	delhi	2020-12-08	
10197	CUST6315	Joshua Martinez	male	25	hyderabad	2022-02-15	
10198	CUST4812	Sarah Miller	female	55	kolkata	2021-03-16	

9000 rows × 12 columns

Correcting Data Types

```
In [41]: df.columns
```

```
Out[41]: Index(['Customer_ID', 'Name', 'Gender', 'Age', 'City', 'Signup_Date',
       'Last_Purchase_Date', 'Purchase_Amount', 'Feedback_Score', 'Email',
       'Phone_Number', 'Country'],
      dtype='object')
```

```
In [42]: df['Age'] = df['Age'].astype('int64')
```

```
In [43]: df['Signup_Date'] = pd.to_datetime(df['Signup_Date'])
```

```
In [44]: df['Last_Purchase_Date'] = pd.to_datetime(df['Last_Purchase_Date'])
```

```
In [45]: df.dtypes
```

```
Out[45]: Customer_ID          object
          Name            object
          Gender           object
          Age             int64
          City            object
          Signup_Date     datetime64[ns]
          Last_Purchase_Date  datetime64[ns]
          Purchase_Amount    float64
          Feedback_Score     float64
          Email            object
          Phone_Number      int64
          Country           object
          dtype: object
```

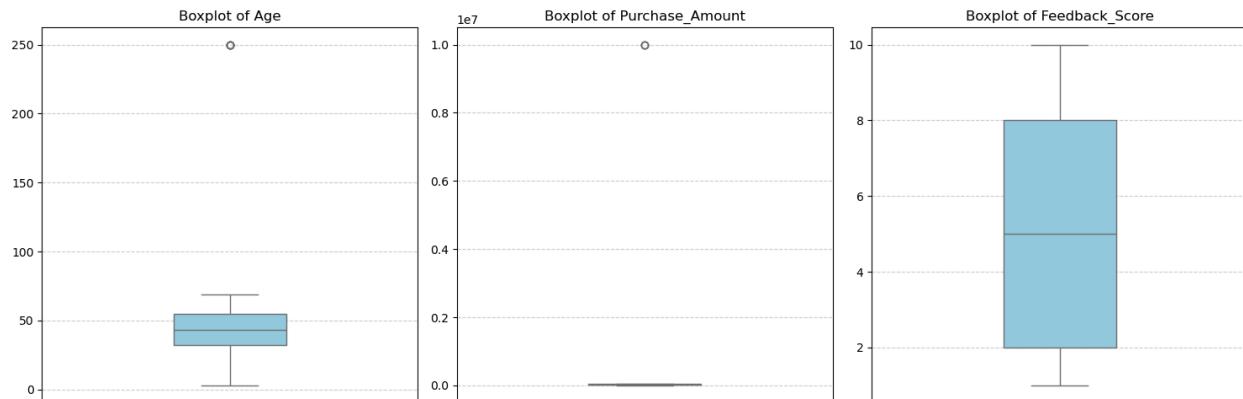
Handling Outliers

```
In [46]: cols = ['Age', 'Purchase_Amount', 'Feedback_Score']

plt.figure(figsize = (15,5))

for i,col in enumerate(cols,1):
    plt.subplot(1,3,i)
    sns.boxplot(y=df[col], color = 'skyblue', width = 0.3)
    plt.title(f'Boxplot of {col}', fontsize = 12)
    plt.ylabel('')
    plt.grid(axis = 'y', linestyle = '--', alpha = 0.6)

plt.tight_layout()
plt.show()
```



```
In [47]: # Calculate Z-scores
z_scores = np.abs(stats.zscore(df[['Age', 'Purchase_Amount']]))


```

```
In [48]: # Identify outliers (any row with z > 3)
df_clean = df[~(z_scores > 3).any(axis = 1)]
```

```
In [49]: cols = ['Age', 'Purchase_Amount', 'Feedback_Score']
```

```
plt.figure(figsize = (15,5))

for i,col in enumerate(cols,1):
    plt.subplot(1,3,i)
    sns.boxplot(y=df_clean[col], color = 'skyblue', width = 0.3)
    plt.title(f'Boxplot of {col}', fontsize = 12)
    plt.ylabel('')
    plt.grid(axis = 'y', linestyle = '--', alpha = 0.6)

plt.tight_layout()
plt.show()
```

